



Vincenzo Innocente: “Optimal floating point computation”

Accuracy, Precision, Speed in scientific computing

- IEEE 754 standard
- Expression optimization
- Approximate Math

- X86_64 SIMD instructions
- Vectorization using the GCC compiler

Objectives

- Understand precision and accuracy in floating point calculations
- Manage optimization, trading precision for speed (or vice-versa!)
- Understand and Exploit vectorization
- Learn how to make the compiler to work for us

Prepare for the exercises

cd hands-on

cd floatingpoint

c++ -O2 -Wall -march=native floats.cpp

./a.out

Try

<https://gcc.godbolt.org/>

Disclaimer, caveats, references

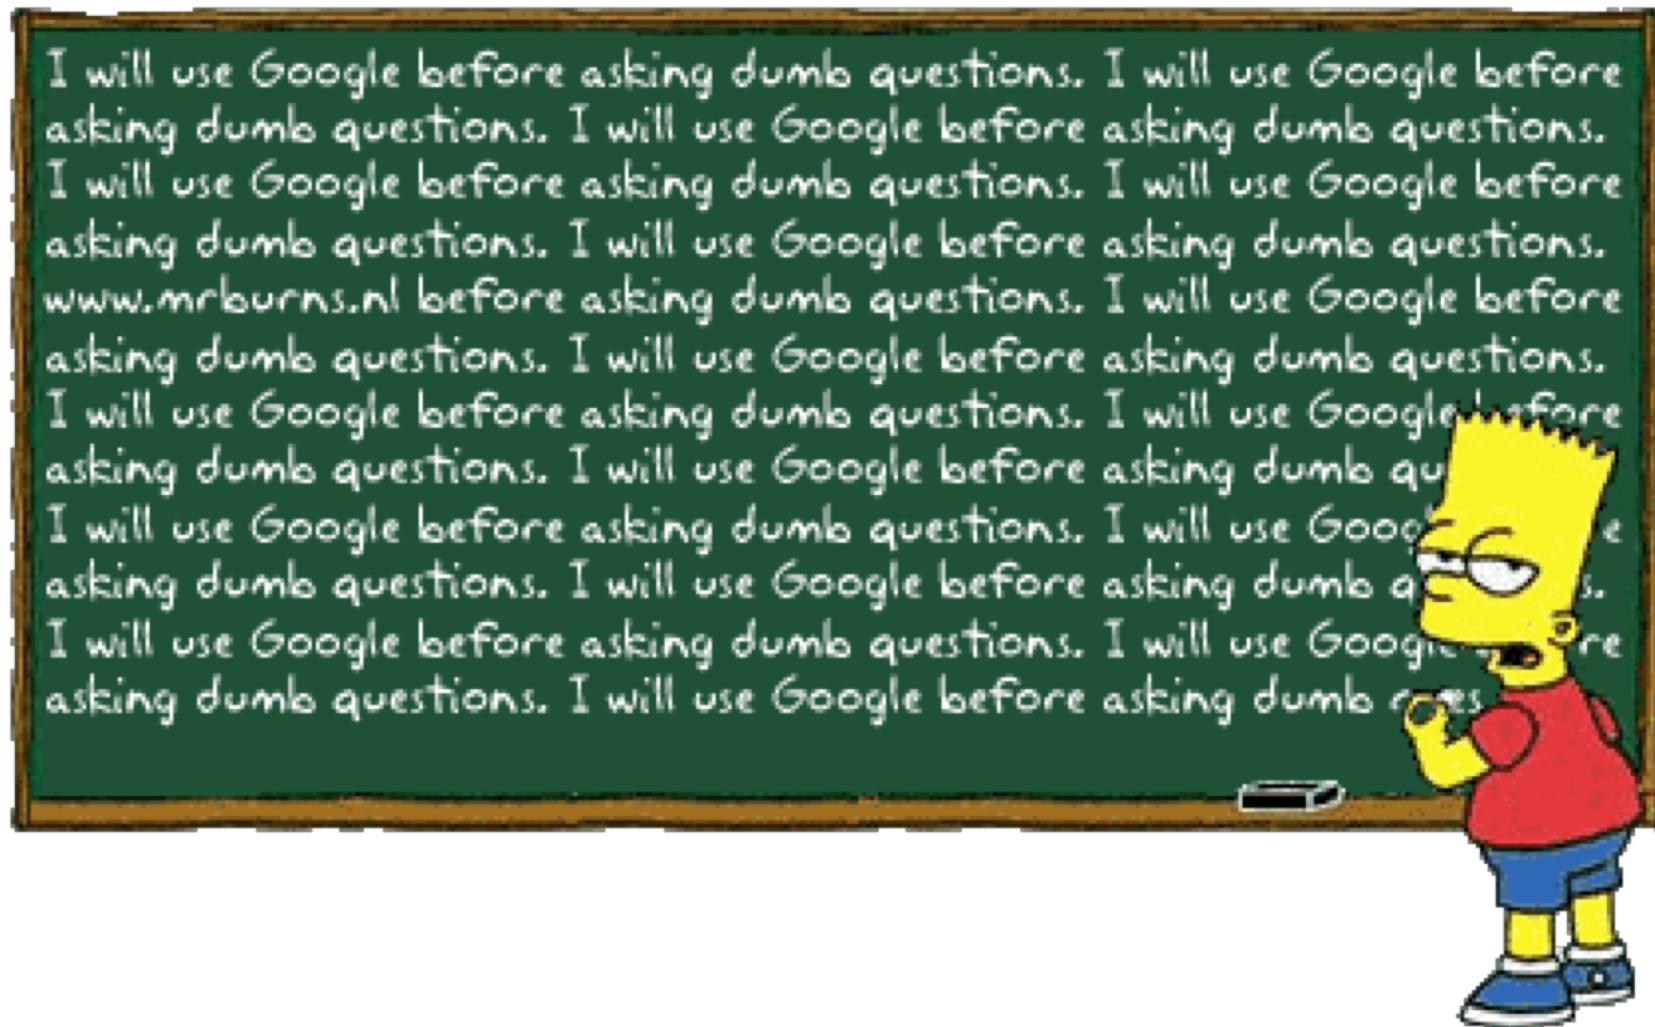
- This is NOT a full overview of IEEE 754
- It applies to x86_64 systems and gcc > 4.7.0
 - Other compilers have similar behavior, details differ
- General References
 - Wikipedia has excellent documentation on the IEEE 754, math algorithms and their computational implementations
 - Handbook of Floating-Point Arithmetic (as google book)
 - Kahan home page
 - INTEL doc and white papers (including ColFax web site)
 - MetaLibm site

Applicability

- It is very (very) different if you deal with
 - Video games
 - Analysis of scientific data
 - Financial applications (with legal bindings)
 - Real time applications
 - Human-life
- <http://www.heidelberg-laureate-forum.org/blog/video/lecture-thursday-september-26-william-morton-kahan/>

Don't be afraid to ask questions!

I will use Google before asking dumb questions. I will use Google before asking dumb questions.

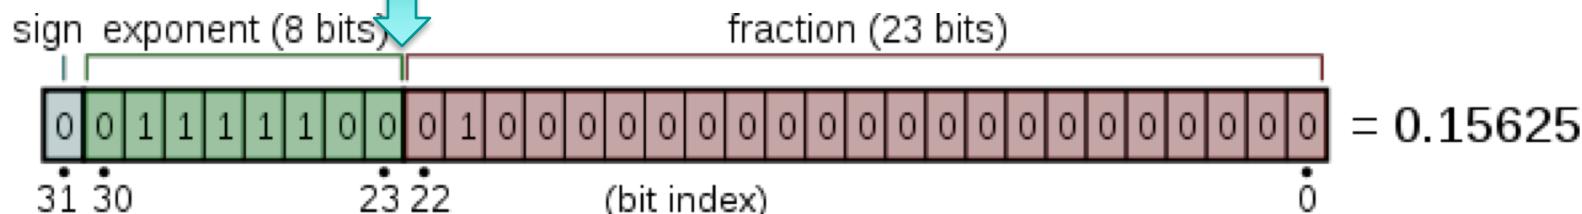


Floating Point Representation

(source Wikipedia)

- floating point describes a system for representing numbers that would be too large or too small to be represented as integers.
 - The advantage of floating-point representation over fixed-point (and integer) representation is that it can support a much wider range of values.
 - int: $-2,147,483,648$ to $+2,147,483,647$, $-(2^{31}) \sim (2^{31}-1)$
 - float: 1.4×10^{-45} to 3.4×10^{38}
 - double: $\sim 10^{-323}$ to $\sim 10^{308}$
 - This has a cost...

IEEE 754 representation of single precision fp



$$n = (-1)^s \times (m2^{-23}) \times 2^{x-127}$$

Exponent	significand zero	significand non-zero	Equation
00 _H	zero, -0	subnormal numbers	$(-1)^{\text{signbits}} \times 2^{-126} \times 0.\text{significandbits}$
01 _H , ..., FE _H	normalized value	normalized value	$(-1^{\text{signbits}} \times 2^{\text{exponentbits}-127} \times 1.\text{significandbits})$
FF _H	±infinity	<u>Nan (quiet,signalling)</u>	

Floating Point Types

Type	Sign	Exponent	Significand field	Total bits	Exponent bias	Bits precision	Number of decimal digits
Half (IEEE 754-2008)	1	5	10	16	15	11	~3.3
Single	1	8	23	32	127	24	~7.2
Double	1	11	52	64	1023	53	~15.9
x86 extended precision	1	15	64	80	16383	64	~19.2
Quad	1	15	112	128	16383	113	~34.0

- Half : `__fp16`, available for storage on recent ARM, INTEL and NVIDIA HW
- Single: `float`, storage and computation in C++ on any hardware
- Double: `double`, storage and computation in C++ on any hardware
- Ext-precision: `long double`, No storage! computation in C++ on x86 hardware
- Quad: `__float128`, storage and (software) computation in C++

C++ interlude

```
#include<iostream>
#include<iomanip>
#include<cmath>
#include<limits>
// std style
template<typename T>
void print(T x) {
    std::cout << std::hexfloat << x << ' ' // exact binary representation
        << std::scientific << std::setprecision(8) << x << ' ' // 8 decimal digits
        << std::defaultfloat << x << std::endl; // 6 decimal digits for float
}

// integer representation
#include<cstring>
#include<bitset>
int f2i(float x) { int i; memcpy(&i,&x,sizeof(int)); return i;} // compiler will optimize
// use bitset to print in binary (bit by bit)
std::cout << std::bitset<32>(f2i(x)) << std::endl; // bit by bit
```

Limits

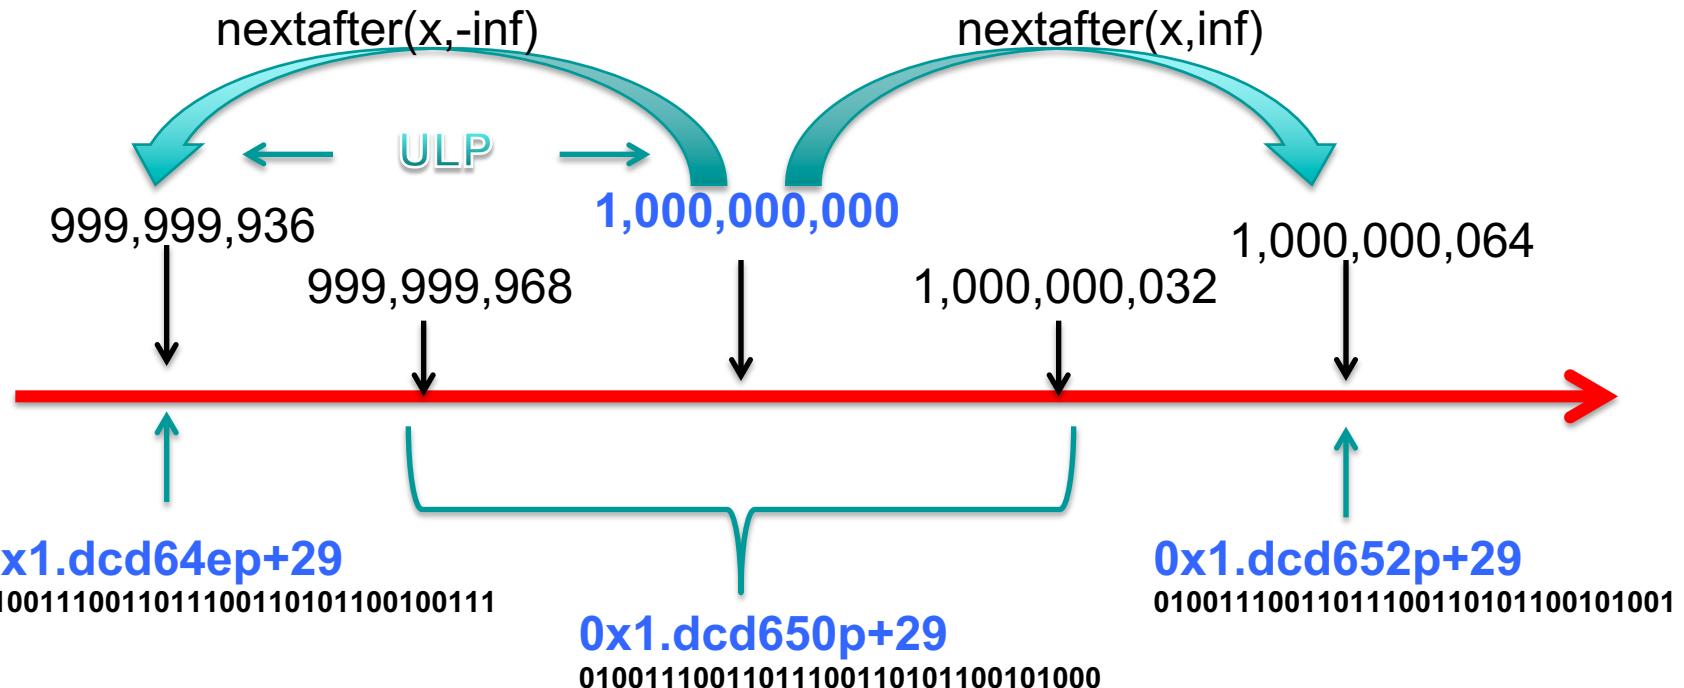
- Everything you want to know about an arithmetic type is in
 - `std::numeric_limits<T>`
 - see http://www.cplusplus.com/reference/limits/numeric_limits/

```
#include <iostream>    // std::cout
#include <limits>      // std::numeric_limits
int main () {
    std::cout << std::boolalpha;
    std::cout << "Minimum value: " << std::numeric_limits<float>::min() << '\n';
    std::cout << "Maximum value: " << std::numeric_limits<float>::max() << '\n';
    std::cout << "Is signed: " << std::numeric_limits<float>::is_signed << '\n';
    std::cout << "significand bits: " << std::numeric_limits<float>::digits << '\n';
    std::cout << "precision: " << std::numeric_limits<float>::epsilon() << '\n';
    std::cout << "has infinity: " << std::numeric_limits<float>::has_infinity << '\n';
    return 0;
}
```

examples

		sign	exponent	significant
0.00000000e+00	0x0p+0	00000000000000000000000000000000		
5.00000000e-01	0x1p-1	00111110000000000000000000000000		
1.00000000e+00	0x1p+0	00111111000000000000000000000000		
1.50000000e+00	0x1.8p+0	00111111100000000000000000000000		
2.00000000e+00	0x1p+1	01000000000000000000000000000000		
2.50000000e+00	0x1.4p+1	01000000010000000000000000000000		
3.00000000e+00	0x1.8p+1	01000000010000000000000000000000		
8.00000000e+00	0x1p+3	01000001000000000000000000000000		
3.14159274e+00	0x1.921fb6p+1	010000001001001000011111011011		
1.0000001e-01	0x1.99999ap-4	00111101110011001100110011001101		
1.17549435e-38	0x1p-126	00000000100000000000000000000000		
3.40282347e+38	0x1.fffffep+127	01111110111111111111111111111111		
nan	nan	11111111100000000000000000000000		
-inf	-inf	11111111100000000000000000000000		
2.93873588e-39	0x1p-128	00000000001000000000000000000000		

Precision, rounding



```
float x = 1000*1000*1000;
```

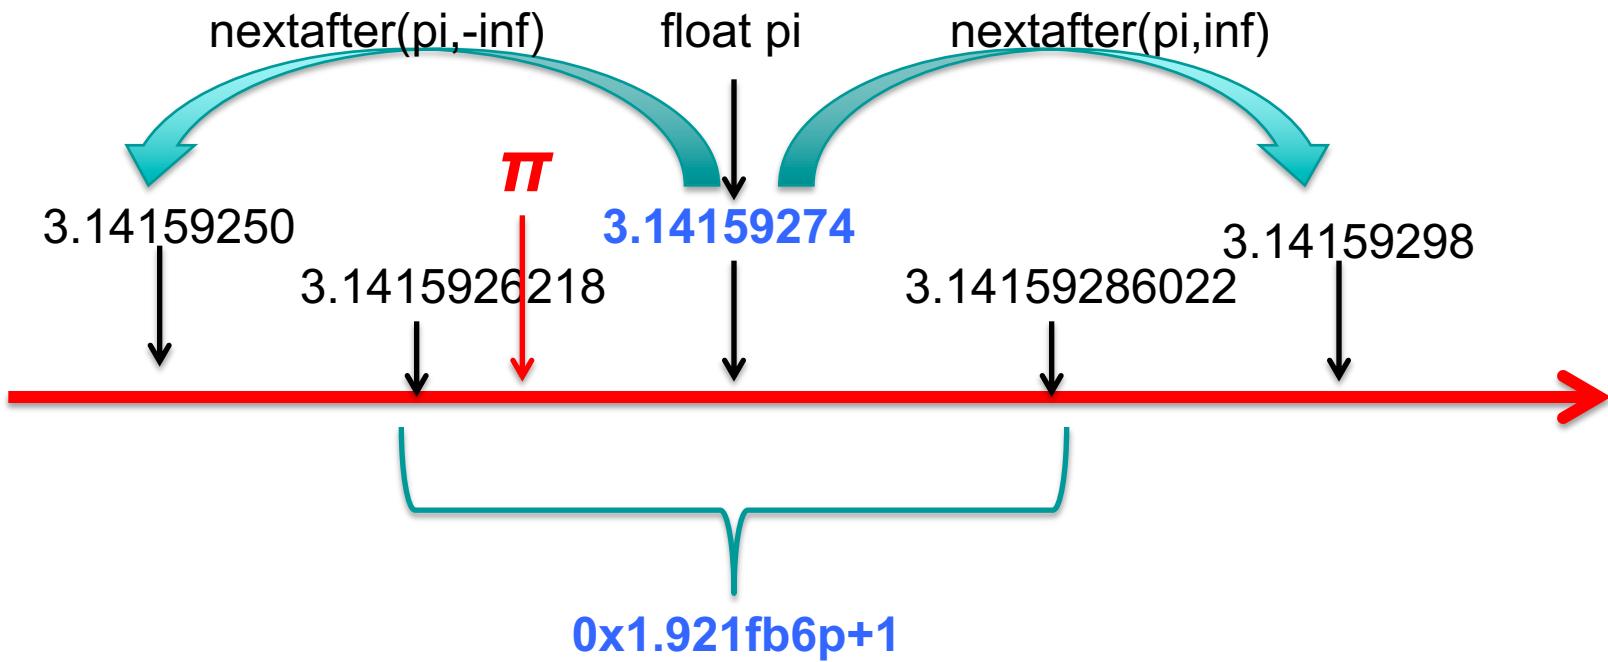
```
std::cout << std::scientific << std::setprecision(8)
    << x << ' ' << x+32.f << ' ' << x+33.f << std::endl
    << x+32.f-x << x+33.f-x << std::endl;
```

1.0000000e+09 1.0000000e+09 1.0000006e+09

0.0000000e+00 6.4000000e+01

https://en.wikipedia.org/wiki/Unit_in_the_last_place

Precision, rounding



```
float pi = M_PI; // 0x1.921fb54442d18p+1 3.1415926535897931e+00
```

Double precision!

Rounding Algorithms

- The standard defines five rounding algorithms.
 - The first two round to a nearest value; the others are called directed roundings:
- Roundings to nearest
 - **Round to nearest, ties to even – rounds to the nearest value;** if the number falls midway it is rounded to the nearest value with an even (zero) least significant bit, which occurs 50% of the time; **this is the default algorithm for binary floating-point** and the recommended default for decimal
 - Round to nearest, ties away from zero – rounds to the nearest value; if the number falls midway it is rounded to the nearest value above (for positive numbers) or below (for negative numbers)
- Directed roundings
 - Round toward 0 – directed rounding towards zero (also known as truncation).
 - Round toward $+\infty$ – directed rounding towards positive infinity (also known as rounding up or ceiling).
 - Round toward $-\infty$ – directed rounding towards negative infinity (also known as rounding down or floor).

Loss of precision

Summing *many* values

```
float x1=0;  
for (float y=1;y<=1000000; ++y) x1+=y;  
print (x1); // 4.99941376e+11 4.9994138e+11 0x1.d19b5p+38
```

```
float x2=0;  
for (float y=1000000;y>0; --y) x2+=y;  
print (x2); // 4.99873677e+11 4.9987368e+11 0x1.d18b2cp+38
```

```
int N=16; float xx[N]; for (int i=0; i<N; ++i) xx[i]=0;  
for (int i=0; i<N; ++i)  
for (int y=i;y<=1000000; y+=N) xx[i]+=float(y);  
float s=0; for (int i=0; i<N; ++i) s+=xx[i];  
print (s); // 4.99999900e+11 4.999999e+11 0x1.d1a944p+38
```

```
float result = 0.5*1000000 *(1000000+1);  
std::cout << std::scientific << (result-x1)/result << " "  
    << (result-x2)/result << " " << (result-s)/result << std::endl;  
// 1.18226832e-04 2.53624079e-04 1.17964680e-06
```

```
float x3=0; float tenth=0.1f;  
for(int i=0; i<10000;++i) x3+=tenth;  
print(x3); // 9.99902893e+02 999.90289 0x1.f3f392p+9 (google for: patriot failure)
```

Loss of precision

Cancellations: Subtracting *too-close* values

For example, consider the quadratic equation:

$$ax^2 + bx + c = 0,$$

with the two exact solutions:

$$x = \frac{-b \pm \sqrt{b^2 - 4ac}}{2a}.$$

Better solution

$$x_1 = \frac{-b - \text{sgn}(b) \sqrt{b^2 - 4ac}}{2a},$$

$$x_2 = \frac{2c}{-b - \text{sgn}(b) \sqrt{b^2 - 4ac}} = \frac{c}{ax_1}.$$

Two sources of cancellation

https://en.wikipedia.org/wiki/Loss_of_significance

Loss of precision

To illustrate the instability of the standard quadratic formula *versus* this variant formula, consider a quadratic equation with roots 1.786737589984535 and $1.149782767465722 \times 10^{-8}$. To sixteen significant figures, roughly corresponding to double-precision accuracy on a computer, the monic quadratic equation with these roots may be written as:

$$x^2 - 1.786737601482363x + 2.054360090947453 \times 10^{-8} = 0$$

Using the standard quadratic formula and maintaining sixteen significant figures at each step, the standard quadratic formula yields

$$\sqrt{\Delta} = 1.786737578486707$$

$$x_1 = (1.786737601482363 + 1.786737578486707)/2 = 1.786737589984535$$

$$x_2 = (1.786737601482363 - 1.786737578486707)/2 = 0.000000011497828$$

Note how cancellation has resulted in x_2 being computed to only eight significant digits of accuracy. The variant formula presented here, however, yields the following:

$$x_1 = (1.786737601482363 + 1.786737578486707)/2 = 1.786737589984535$$

$$x_2 = 2.054360090947453 \times 10^{-8} / 1.786737589984535 = 1.149782767465722 \times 10^{-8}$$

Note the retention of all significant digits for x_2 .

Floating Point Math

- Floating point numbers are NOT real numbers
 - They exist in a finite number ($\sim 2^{32}$ for single prec)
 - Exercise: count how many floats exists between given two
 - Exist a “next” and a “previous” (std::nextafter)
 - Differ of one ULP (Unit in the Last Place or Unit of Least Precision
http://en.wikipedia.org/wiki/Unit_in_the_last_place)
 - Results of Operations are rounded
 - Standard conformance requires half-ULP precision.
 - $x + \varepsilon - x \neq \varepsilon$ (can easily be 0 or ∞)
 - Their algebra is not associative
 - $(a+b) + c \neq a+(b+c)$
 - $a/b \neq a*(1/b)$
 - $(a+b)*(a-b) \neq a^2 - b^2$

Strict IEEE754 vs “Finite” (fast) Math

- Compilers can treat FP math either in “strict IEEE754 mode” or optimize operations using “algebra rules for finite real numbers” (as in FORTRAN)
 - gcc default is “strict IEEE754 mode”
 - `-O2 -funsafe-math -ffast-math`
 - `-Ofast` (switch vectorization on as well)
 - Caveat: the compiler improves continuously: still it does not optimize yet all kind of expressions
- <https://gcc.gnu.org/wiki/FloatingPointMath>

Inspecting generated code

`objdump -S -r -C --no-show-raw-instr -w kernel.o | less`
(on MacOS: `otool -t -v -V -X kernel.o | c++filt | less`)

Or use <http://gcc.godbolt.org>

c++ -S kernel.cc; less kernel.s

```
float
kernel(float a, float x, float y)
{
    return a*x + a*y;
}
```

-O2

kernel(float, float, float):

mulss	%xmm0,%xmm2
mulss	%xmm0,%xmm1
addss	%xmm2,%xmm1
movaps	%xmm1,%xmm0
ret	

-Ofast

kernel(float, float, float):

addss	%xmm2,%xmm1
mulss	%xmm0,%xmm1
movaps	%xmm1,%xmm0
ret	

Exercise: compare assembler for O2 and Ofast
<http://goo.gl/5jwOVO>

FMA (fused multiply add)

- $fma(a,b,c) = \text{round}(a*b+c)$
 - Opposed to $\text{round}(\text{round}(a*b)+c)$
- Single instruction with 4 or 5 cycle latency
 - Opposed to 2 instructions with 5+3 cycle latency
- More precise (one rounding instead of two)
 - Results will differ!
- Introduce a “contraction” issue
 - No standard contraction rules (yet)
 - $a*b+c*d$, $a+b*c+d$, $x^*=a;x+=b$, etc

Controlling fma (contraction)

■ gcc (and clang) flags

- `-ffp-contract=off` (default is “fast”)
- <https://stackoverflow.com/questions/43352510/difference-in-gcc-ffp-contract-options>

■ cuda (nvcc) flags

- `-fmad=false` (default is “true”)
- <https://docs.nvidia.com/cuda/floating-point/index.html>

fma library function

- **std::fma**
 - Implemented in software (libm)
 - Used if fma instruction not supported: SLOW
 - Compilers will use fma instruction if exists
 - Not affected by re-association flags
- To fully control fma and be fully reproducible
 - use std::fma
 - disable “contraction”

Gradual underflow (subnormals)

- *Subnormals* (or *denormals*) are fp smaller than the smallest normalized fp: they have leading zeros in the significand
 - For single precision they represent the range 10^{-38} to 10^{-45}
- Subnormals guarantee that additions never underflow
 - Any other operation producing a *subnormal* will raise a underflow exception if also inexact
- Literature is full of very good reasons why “gradual underflow” improves accuracy
 - This is why they are part of the IEEE 754 standard
- Hardware is not always able to deal with *subnormals*
 - Software assist is required: **SLOW**
 - To get correct results even the software algorithms need to be specialized
- It is possible to tell the hardware to *flush-to-zero* subnormals
 - It will raise underflow and inexact exceptions

Floating point exceptions

The IEEE floating point standard defines several exceptions that occur when the result of a floating point operation is unclear or undesirable. Exceptions can be ignored, in which case some default action is taken, such as returning a special value. When trapping is enabled for an exception, an error is signalled whenever that exception occurs. These are the possible floating point exceptions:

- **Underflow:** This exception occurs when the result of an operation is too small to be represented as a normalized float in its format. If trapping is enabled, the *floating-point-underflow* condition is signalled. Otherwise, the operation results in a denormalized float or zero.
- **Overflow:** This exception occurs when the result of an operation is too large to be represented as a float in its format. If trapping is enabled, the *floating-point-overflow* exception is signalled. Otherwise, the operation results in the appropriate infinity.
- **Divide-by-zero:** This exception occurs when a float is divided by zero. If trapping is enabled, the *divide-by-zero* condition is signalled. Otherwise, the appropriate infinity is returned.
- **Invalid:** This exception occurs when the result of an operation is ill-defined, such as $(0.0/0.0)$. If trapping is enabled, the *floating-point-invalid* condition is signalled. Otherwise, a quiet NaN is returned.
- **Inexact:** This exception occurs when the result of a floating point operation is not exact, i.e. the result was rounded. If trapping is enabled, the *floating-point-inexact* condition is signalled. Otherwise, the rounded result is returned.

Defensive Computing

- Verify pre/post conditions and algorithm invariants
- Strict IEEE754 conformance
- In case of invalid input notify
 - Exception
 - Error code
 - Nan, inf..
- Can be implemented as a wrapper of a HPC algo

High Performance

- Garbage in, garbage out
- Assume finite math
- Invalid input
 - Undefined behavior
 - Algorithm processes it at no additional cost
- User responsibility to guarantee correctness

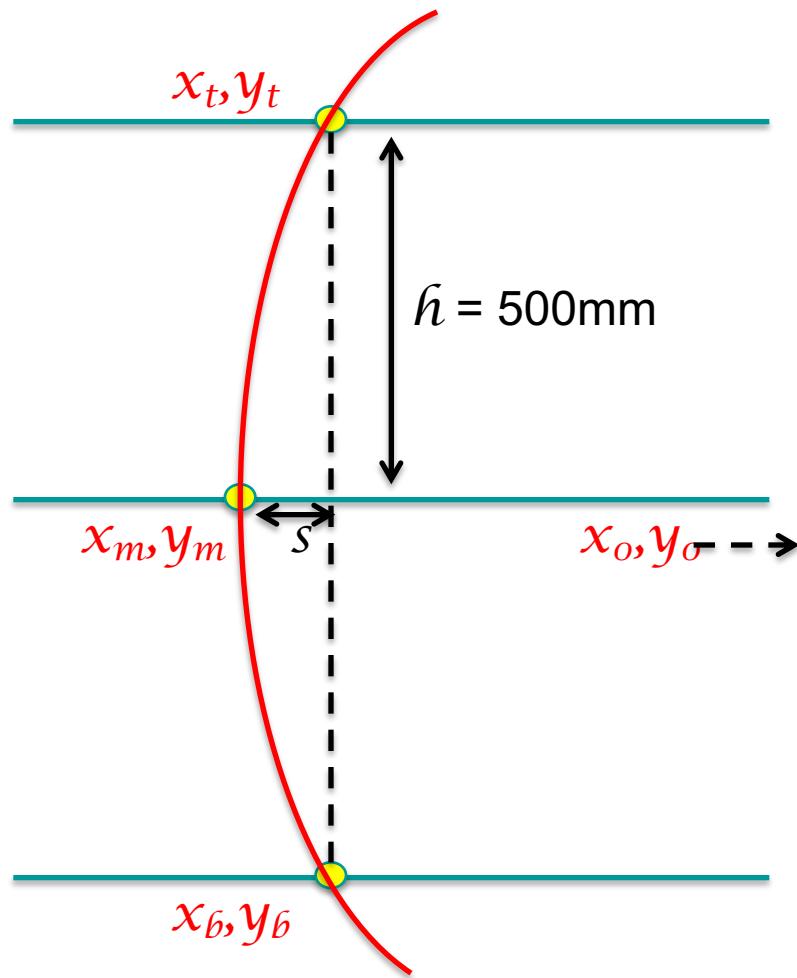
In <https://gcc.godbolt.org/z/OmNMxx> change -O3 in -Ofast

PRECISION, ACCURACY, SPEED

Definitions (mine)

- Precision
 - Numerical precision (in bits)
 - Can be evaluated using a higher “precision”
- Accuracy
 - The error w/r/t the truth
 - Can be evaluated comparing different algorithms
- Target Accuracy
 - The acceptable tolerance w/r/t the above
 - Evaluated, for instance, w/r/t known error on the truth

Example: detector of size 1 m and resolution 10 micron (required accuracy: ~ 1 micron?)



three measurement planes

Is single precision enough to represent

- geometric elements?
- Measurements?
- computed trajectory?

Is enough to perform computation?

Trajectory: circle with radius r
 s is the sagita, h the half-chord

$$r = (s^2 + h^2) / 2s \sim h^2 / 2s$$

$$h^2 = 2sr - s^2$$

$$x_o = r + x_m$$

Circle equation and solutions

$$(x - \textcolor{red}{x}_o)^2 + (y - \textcolor{red}{y}_o)^2 = r^2$$

Compute $\textcolor{red}{x}_t$ (x for $y=h$)

$$x = \textcolor{red}{x}_o - (r^2 - h^2)^{1/2} \sim o \text{ (zero)}$$

Alternative formulation equivalent to change of coordinate centered at $\textcolor{red}{x}_m, \textcolor{red}{y}_m$

$$(x - \textcolor{red}{x}_m)^2 + (y - \textcolor{red}{y}_m)^2 - 2ar(x - \textcolor{red}{x}_m) - 2br(y - \textcolor{red}{y}_m) = o$$

In our case $a=1, b=0$

$$(x - \textcolor{red}{x}_m)^2 - 2(x - \textcolor{red}{x}_m)r + h^2 = o$$

$$x = \textcolor{red}{x}_m + h^2 / (r + (r^2 - h^2)^{1/2}) \sim \textcolor{red}{x}_m + h^2 / 2r$$

A correct choice of algorithm can allow to use a reduced precision to obtain the same accuracy

- Faster
- Smaller memory footprint

PRECISION, ACCURACY, SPEED

Cost of operations (in cpu cycles)

op	instruction	sse s	sse d	avx s	avx d (FMA s&d)	
+,-	ADD,SUB	3	3	3	3	4
== < >	COMISS CMP..	2,3	2,3	2,3	2,3	
f=d d=f	CVT..	3	3	4	4	
,&,^	AND,OR	1	1	1	1	
*	MUL	5	5	5	5	4
/,sqrt	DIV, SQRT	10-14	10-22	21-29	21-45	
1.f/ , 1.f/sqrt	RCP, RSQRT	5		7		
=	MOV	1,3,...	1,3,...	1,4,....	1,4,...	→350 from main memory

Approximate reciprocal (sqrt, div)

The major contribution of game industry to SE is the discovery of the “magic” fast $1/\sqrt{x}$ algorithm

```
float invSqrt(float x){  
    int i; memcpy(&i,&x,4);  
    i = 0x5f3759df - (i >> 1);  
    float y; memcpy(&y,&i,4); // approximate  
    return y * (1.5f - 0.5f * x * y * y); // better  
}
```

Compilers use them at Ofast
Accuracy 2ULP
Non standard: Hardware dependent

Real x86_64 code for $1/\sqrt{x}$

```
_mm_store_ss( &y, _mm_rsqrt_ss( _mm_load_ss( &x ) ) );  
return y * (1.5f - 0.5f * x * y * y); // One round of Newton's method
```

Real x86_64 code for $1/x$

```
_mm_store_ss( &y, _mm_rcp_ss( _mm_load_ss( &x ) ) );  
return (y+y) - x*y*y; // One round of Newton's method
```

Cost of functions (in cpu cycles i7sb)

	Gnu libm (~IEEE754 compliant)		VDT (Cephes) (~1ULP precise)	
	single	double	single	double
sin,cos,tan large x	55 >500	100	30	50
sincos	70		40	
atan2	50	100	30	45
exp	650	65	42	55
log	50	105	37	42

`SET_RESTORE_ROUND_NOEXF (FE_TONEAREST);`

Precision vs Speed in “libm”

■ Log

- 24 bit precision needs 8th degree polynomial
- 16 bit precision needs 5th degree polynomial
- 7 bit precision needs 2th degree polynomial

	2	3	4	5	6	7	8
e^x	6.8	10.8	15.1	19.8	24.6	29.6	34.7
$\sin(x)$	7.8	12.7	16.1	21.6	25.5	31.3	35.7
$\ln(1+x)$	8.2	11.1	14.0	16.8	19.6	22.3	25.0
$\arctan(x)$	8.7	9.8	13.2	15.5	17.2	21.2	22.3
$\tan(x)$	4.8	6.9	8.9	10.9	12.9	14.9	16.9
$\arcsin(x)$	3.4	4.0	4.4	4.7	4.9	5.1	5.3
\sqrt{x}	3.9	4.4	4.8	5.2	5.4	5.6	5.8

Number of significant bits of accuracy vs degree of minimax polynomial approximating the range [0..1].

Example: multiple scattering formula

```
double ms(double radLen, double m2, double p2) {  
    constexpr double amscon = 1.8496e-4; // (13.6MeV)**2  
    double e2 = p2 + m2;  
    double beta2 = p2/e2;  
    double fact = 1 + 0.038*log(radLen); fact *=fact;  
    double a = fact/(beta2*p2);  
    return amscon*radLen*a;  
}
```

Already an approximation

```
float msf(float radLen, float m2, float p2) {  
    constexpr float amscon = 1.8496e-4; // (13.6MeV)**2  
    float e2 = p2 + m2;
```

```
    float fact = 1.f + 0.038f*unsafe_logf<2>(radLen); fact /= p2;  
    fact *=fact;  
    float a = e2*fact;  
    return amscon*radLen*a;  
}
```

Material density,
thickness, track angle
Known at percent?

2nd order polynomial

How to speed up math

- Avoid or factorize-out division and sqrt
 - if possible compile with “–Ofast”
- Prefer linear algebra to trigonometric functions
- Cache quantities often used
 - No free lunch: at best trading memory for cpu
- Choose precision to match required accuracy
 - Square and square-root decrease precision
 - Catastrophic precision-loss in the subtraction of almost-equal large numbers

Example: cut in pt, phi, eta

```
inline float pt2(float x, float y) {return x*x+y*y;}
inline float pt(float x, float y) {return std::sqrt(pt2(x,y));}
inline float phi(float x, float y) {return std::atan2(y,x);}
inline float eta(float x, float y, float z) { float t(z/pt(x,y)); return ::asinhf(t);}
inline float dot(float x1, float y1, float x2, float y2) {return x1*x2+y1*y2;}
inline float dphi(float p1,float p2) {
    auto dp=std::abs(p1-p2); if (dp>float(M_PI)) dp-=float(2*M_PI);
    return std::abs(dp);
};
```

```
if (pt(x[i],y[i])>ptcut) ...
if (dphi(phi(x[i],y[i]),phi(x[j],y[j]))<phicut) ....
```

Exercise in ptcut.cpp

Formula Translation: what was the author's intention?

```
// Energy loss and variance according to Bethe and Heitler, see also  
// Comp. Phys. Comm. 79 (1994) 157.  
//  
double p = mom.mag();  
double normalisedPath = fabs(p/mom.z())*radLen;  
double z = exp(-normalisedPath);  
double varz = (exp(-normalisedPath*log(3.)/log(2.))- exp(-2*normalisedPath));
```

```
double pt = mom.perp();  
double j = a_i*(d_x * mom.x() + d_y * mom.y())/(pt*pt);  
double r_x = d_x - 2* mom.x()*(d_x*mom.x()+d_y*mom.y())/(pt*pt);  
double r_y = d_y - 2* mom.y()*(d_x*mom.x()+d_y*mom.y())/(pt*pt);  
double s = 1/(pt*pt*sqrt(1 - j*j));
```

Exercise: edit Optimizelt.cc to make it “optimal”
check the generated assembly, verify speed gain

Formula Translation: the solution?

Is the compiler able to do it for us?

SIMD: LOOP VECTORIZATION

Test environment

```
cd hands-on/vectorization
```

```
c++ -O2 pi.cpp -fopt-info-vec -march=native;  
time ./a.out
```

```
c++ -Ofast pi.cpp -fopt-info-vec -march=native;  
time ./a.out
```

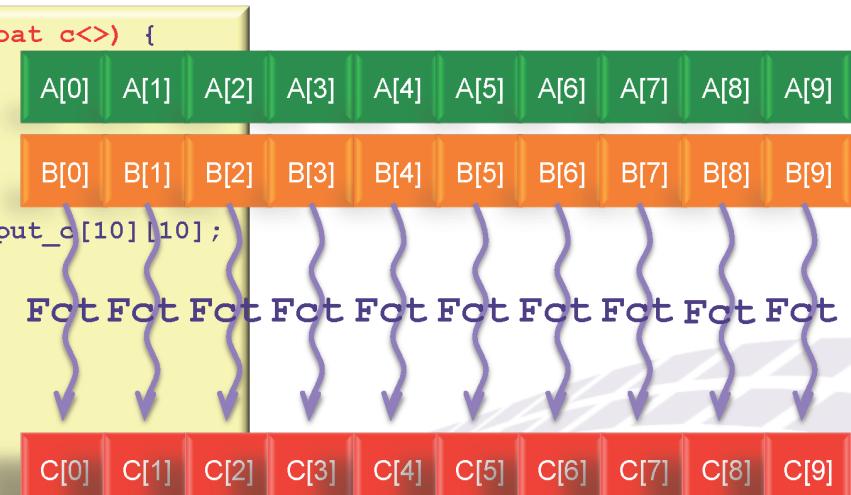
What is Stream Computing?



- A similar computation is performed on a collection of data (*stream*)
 - There is no data dependence between the computation on different stream elements
- Stream programming is well suited to GPU *and vector-cpu!*

```
kernel void Fct(float a<>, float b<>, out float c<>) {
    c = a + b;
}
int main(int argc, char** argv) {
    int i, j;
    float a<10, 10>, b<10, 10>, c<10, 10>;
    float input_a[10][10], input_b[10][10], input_c[10][10];
    for(i=0; i<10; i++) {
        for(j=0; j<10; j++) {
            input_a[i][j] = (float) i;
            input_b[i][j] = (float) j;
        }
    }
    streamRead(a, input_a);
    streamRead(b, input_b);
    Fct(a, b, c);
    streamWrite(c, input_c);
    ...
}
```

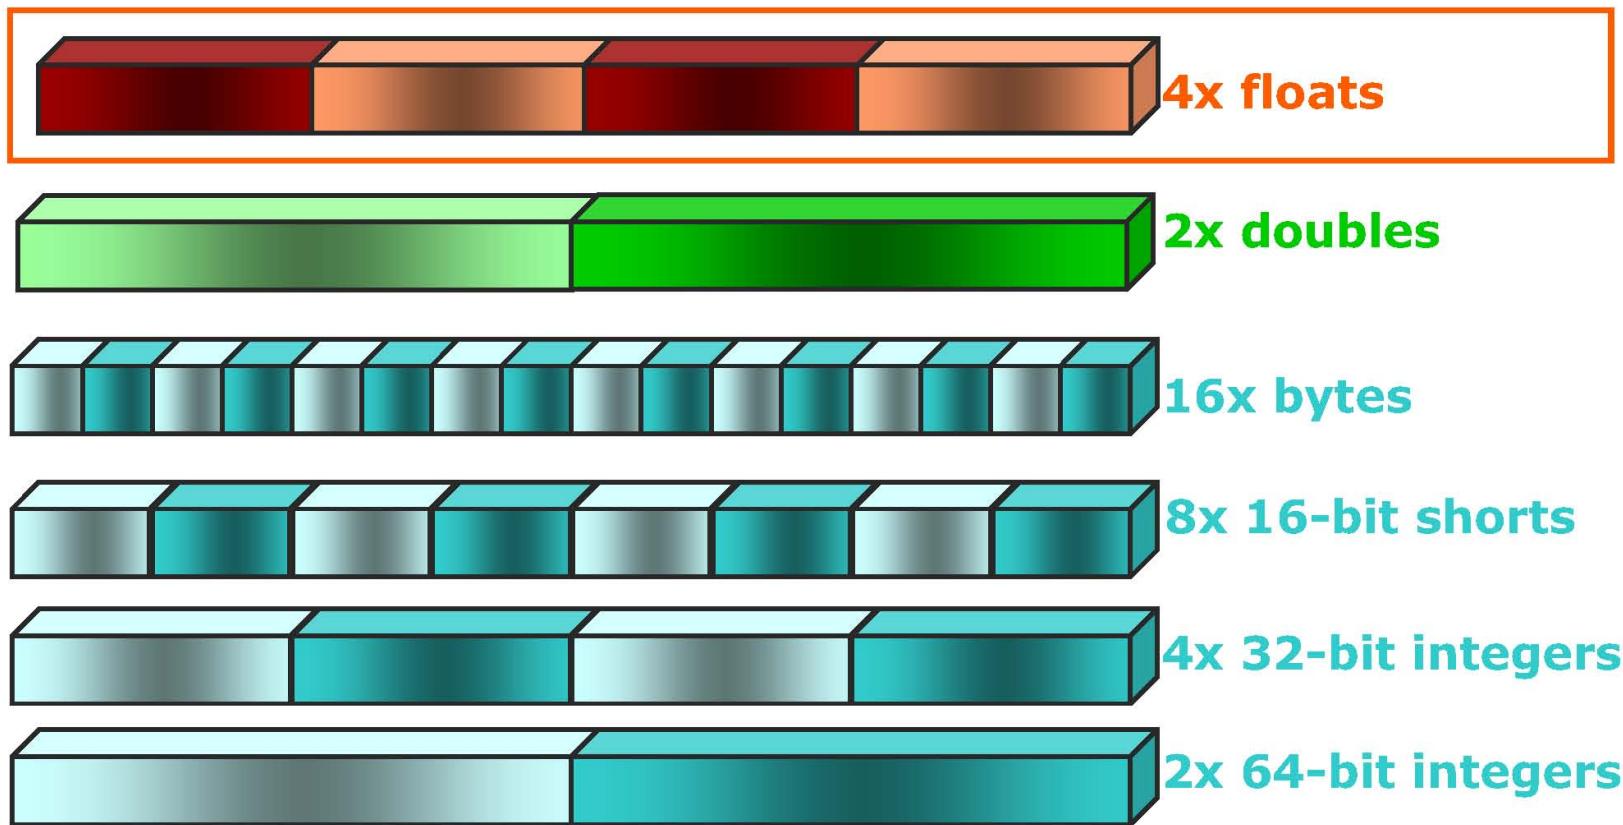
Brook+ example



www.caps-enterprise.com

SSE Data types

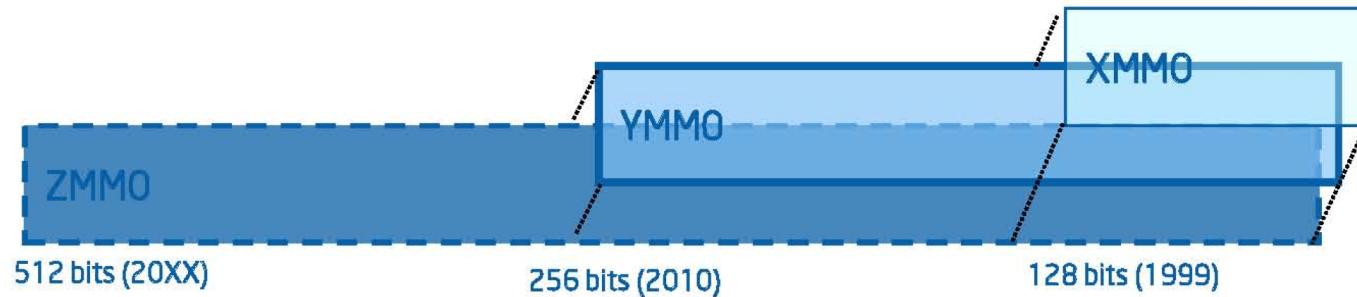
128 bit word (XMM)

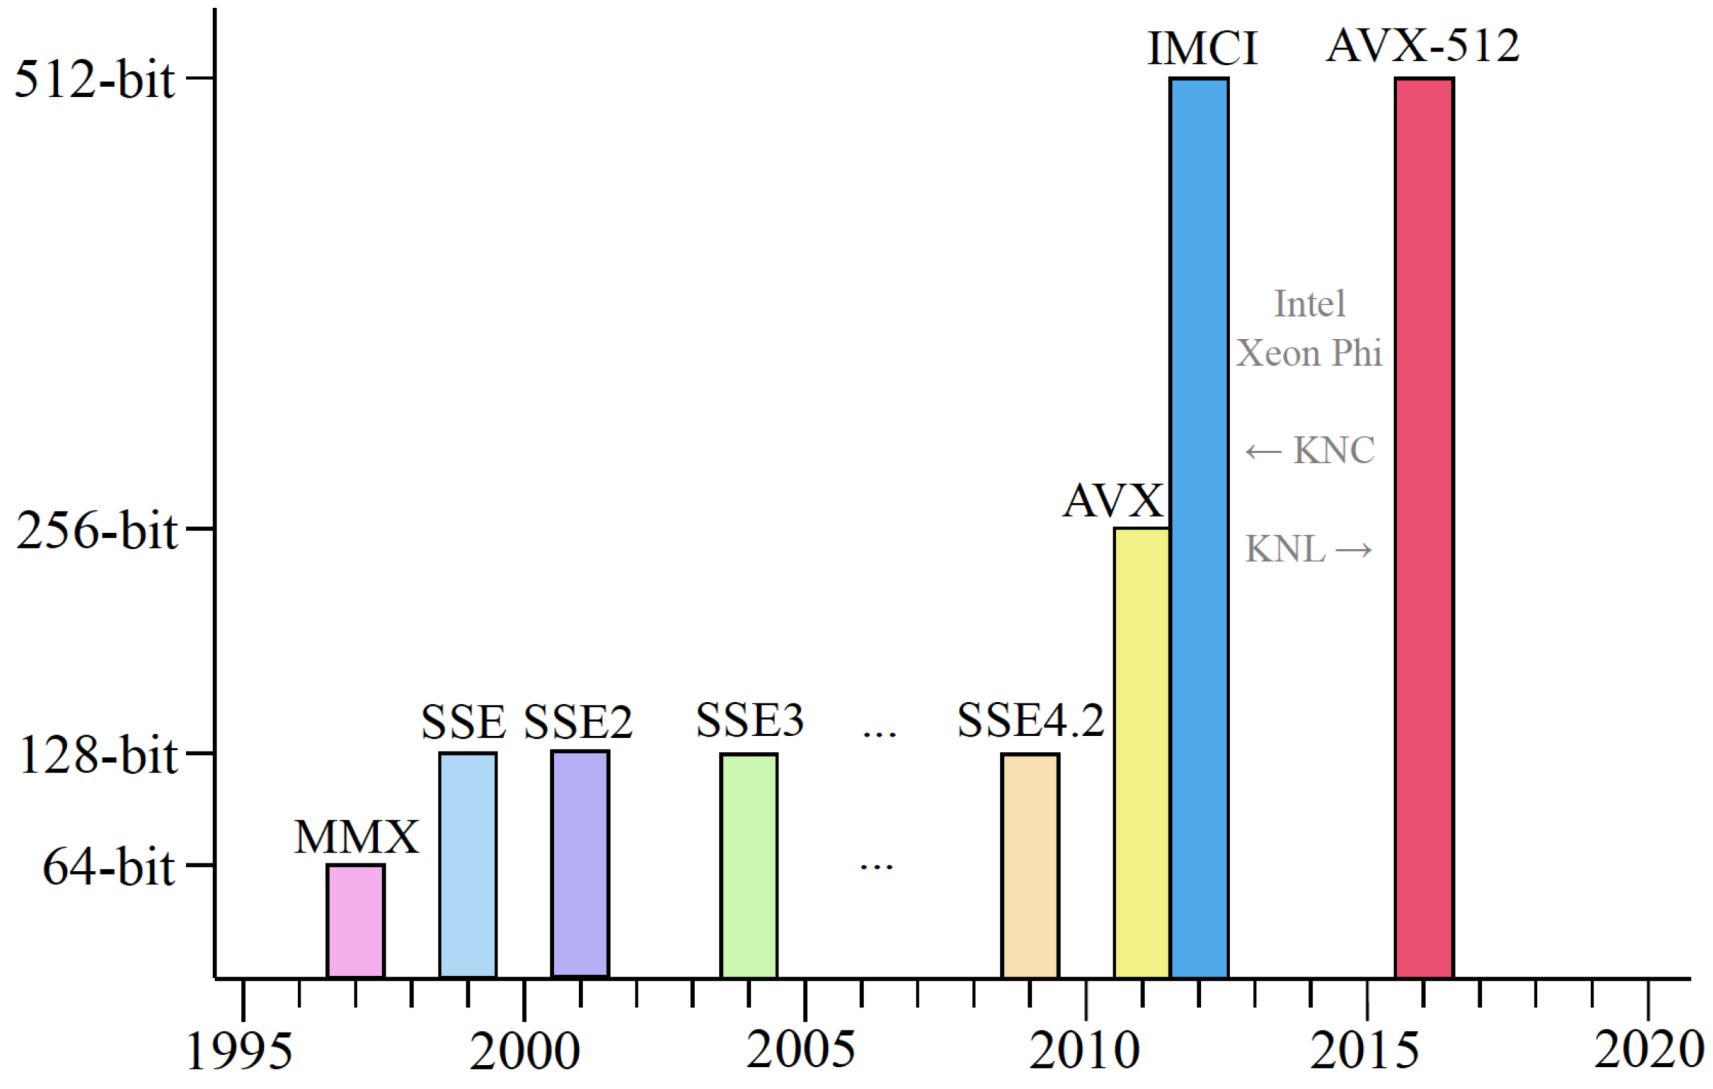


Sandy Bridge (2010): Intel® AVX

A 256-bit vector extension to SSE

- Intel® AVX extends all 16 XMM registers to 256bits
- Intel® AVX works on either
 - The whole 256-bits
 - The lower 128-bits(like existing SSE instructions)
 - A drop-in replacement for all existing scalar/128-bit SSE instructions
- The new state extends/overlays SSE
- The lower part (bits 0-127) of the YMM registers is mapped onto XMM registers





Single Instruction Multiple Data

- **Scalar processing**

- traditional mode
- one operation produces one result



X

x3

x2

x1

x0

+

Y

y3

y2

y1

y0

X + Y

x3+y3

x2+y2

x1+y1

x0+y0

Slide Source: Alex Klimovitski & Dean Macri, Intel Corporation

Single Instruction Multiple Data

- SLP (Superword Level Parallelism)
 - Direct mapping to underling SIMD machine instruction
 - Usually implemented using array/vector notation
- Loop Vectorization
 - Transform a loop into N streams ($N = \text{SIMD-width}$)
 - Compiler assisted or implemented in a “vector-library”
- Loop vectorization is more efficient than SLP
 - Transform your problem in a long loop over simple quantities

“Vector Extension”

- SIMD vector can be found implemented as compiler's extension or libraries
 - <https://gcc.gnu.org/onlinedocs/gcc/Vector-Extensions.html>
 - <http://clang.llvm.org/docs/LanguageExtensions.html#vectors-and-extended-vectors>
 - <http://www.cilkplus.org/tutorial-array-notation>
 - <http://code.compeng.uni-frankfurt.de/projects/vc>
- Here we will experiment with “gcc vectors”

Gcc Vector extension

```
typedef T __attribute__( ( vector_size( N*sizeof(T) ) ) ) VecNT;
```

```
typedef float __attribute__( ( vector_size( 16 ) ) ) Vec4F;
```

```
Vec4F a{0,1f,-2f,3f}, b{-1f,-2f,3f,0}, c{0,2.f,4.f,5.f}, zero{0};
```

```
c += 3.14f*b*a;  
auto z = (a>0) ? a : -a;  
auto m = (a>b) ? a : b;  
auto t = a/b; t = (x>pi/8.f) ? (t-1.0f)/(t+1.0f) : t;  
❑ scalar code: if(x>pi/8.f) t= (t-1.0f)/(t+1.0f) ;
```

**Vector of integers:
0 for false, -1 for true**

always computed

“Auto” vectorization

- The process that transform scalar code in vector code
- As a compiler pass
 - Issues: detect and manage data dependencies
 - Hints from user as macro
- OpenMP4
 - New: implementations still experimental
 - icc: “fully supported?”
 - gcc: syntax and code generation ok, vectorization relies on the corresponding compiler pass
- OpenCL:
 - Uses compiler backend

Loop Vectorization

More at <http://gcc.gnu.org/projects/tree-ssa/vectorization.html>

- ❖ original serial loop:

```
for(i=0; i<N; ++i){  
    a[i] = a[i] + b[i];  
}
```

vectorization

- ❖ loop in vector notation:

```
for (i=0; i<N/4; ; i+=4){  
    V4& va = *((V4*)a[i]);  
    va += *((V4 const *)b[i]);  
}
```

vectorized loop (4 times less iterations)

- ❖ loop in vector notation:

```
for (i=0; i<(N-N%4); i+=4){  
    V4& va = *((V4*)a[i]);  
    va += *((V4 const *)b[i]);  
}
```

vectorized loop

```
for ( ; i < N; i++) {  
    a[i] = a[i] + b[i];  
}
```

epilog loop

- ❖ Loop based vectorization

- ❖ No dependences between iterations

Loop Dependencies

```
for (i=0; i<N; ++i){  
    A[i+1] = B[i] + X  
    D[i] = A[i] + Y  
}
```

```
for (i=0; i<N; ++i){  
    B[i] = A[i] + Y  
    A[i+1] = B[i] + X  
}
```

```
for (i=0; i<N, ++i)  
    for (j=0, j<N; ++j)  
        A[i+1][j] = A[i][j] + X
```

```
for (i=0; i<N; ++i)  
    A[i+1] = B[i] + X
```

```
for (i=0; i<N; ++i)  
    D[i] = A[i] + Y
```

```
for (i=0; i<N; ++i){  
    B[i] = A[i] + Y  
    A[i+1] = B[i] + X  
}
```

Subtle issues:

A may partially overlap
with B
with X
with N

The compiler will generate
runtime checks and alternative
sequential code.

To avoid this overhead
use local variable,
“restrict” keyword or
#pragma GCC ivdep

```
void ignore_vec_dep (int *a, int k, int c, int m)  
{  
#pragma GCC ivdep  
    for (int i = 0; i < m; ++i)  
        a[i] = a[i + k] * c;  
}
```

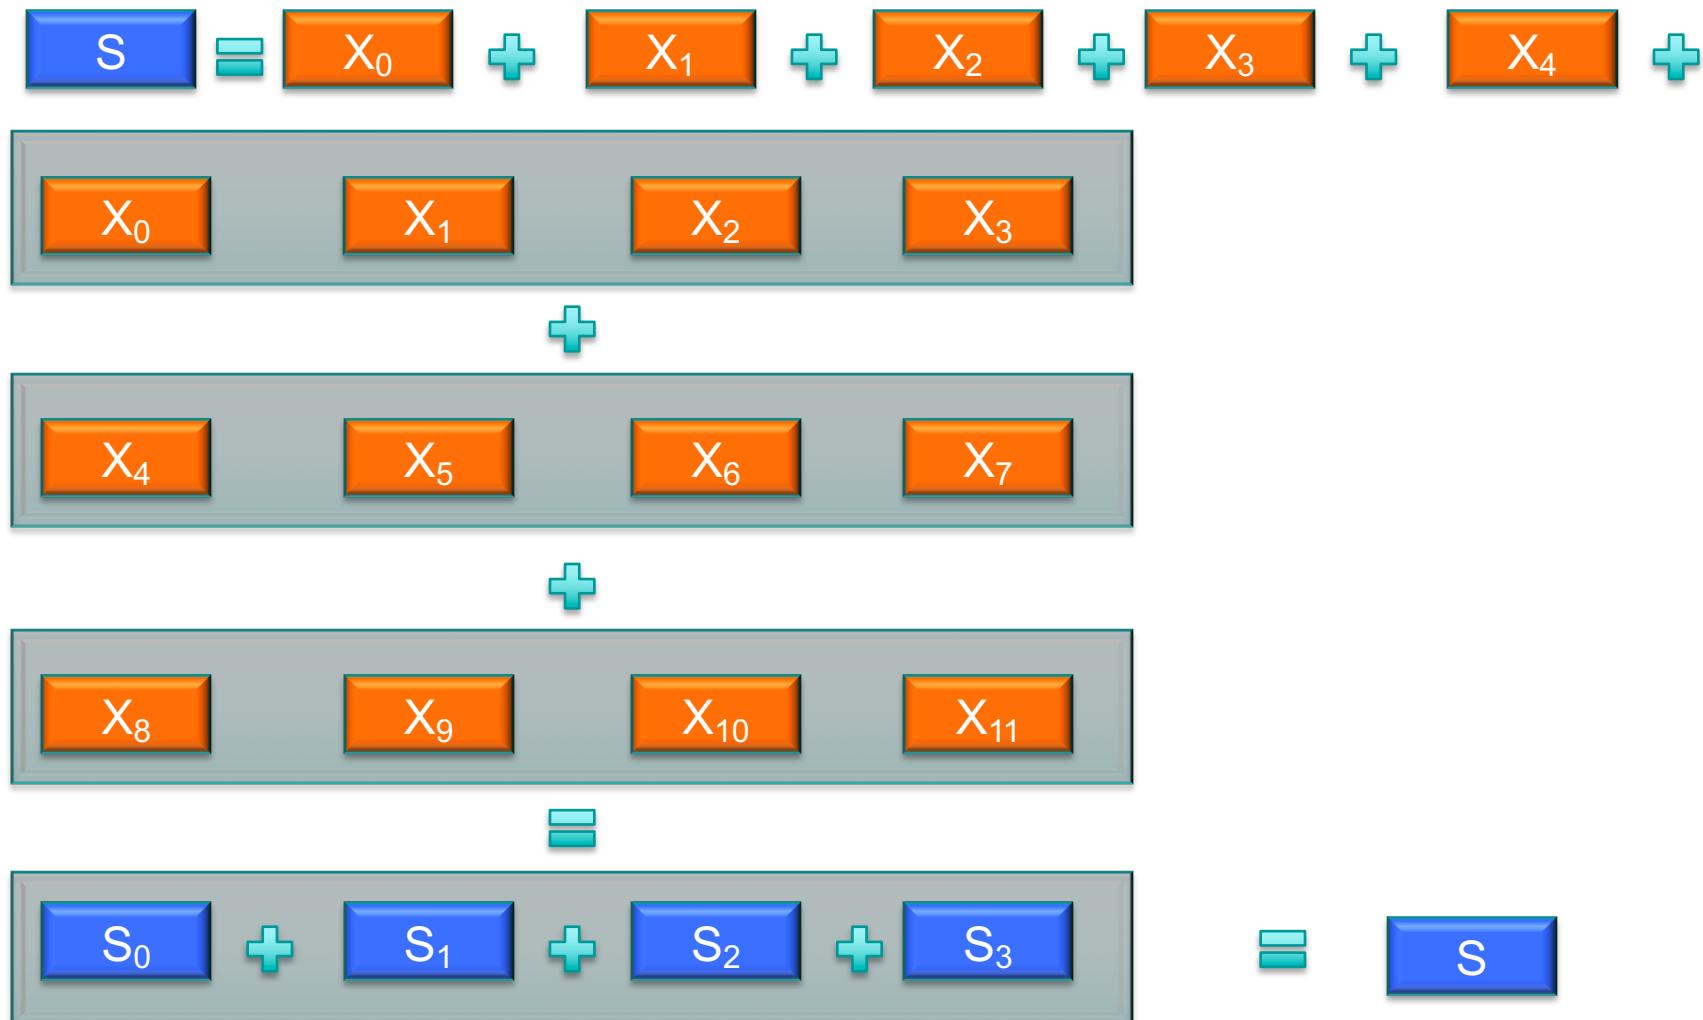
Reduction

```
float innerProduct() {  
    float s=0;  
    for (int i=0; i!=N; i++)  
        s+= a[i]*b[i];  
    return s;  
}
```

```
// pseudo code  
float innerProduct() {  
    V4 s ={0};  
    for (int i=0; i<N/4; ++i)  
        s += *((V4 const *)a[i]) * *((V4 const *)b[i]);  
  
    // horizontal sum;  
    float sum=s[0];  
    for (int j=1;j!=4;++j)sum+=s[j];  
    return sum;  
}
```

Requires *relaxed* float-math rules
(-Ofast)

Reduction



Real code <http://goo.gl/3ku9EJ>

```
float a[1024],b[1024]; int N=1024;  
float innerProduct() {  
    float s=0;  
    for (int i=0; i!=N; i++)  
        s+= a[i]*b[i];  
    return s;  
}
```

.L3:

```
    movss  a(%rax), %xmm1  
    addq   $4, %rax  
    mulss  b-4(%rax), %xmm1  
    addss  %xmm1, %xmm0  
    cmpq   %rdx, %rax  
    jne    .L3
```

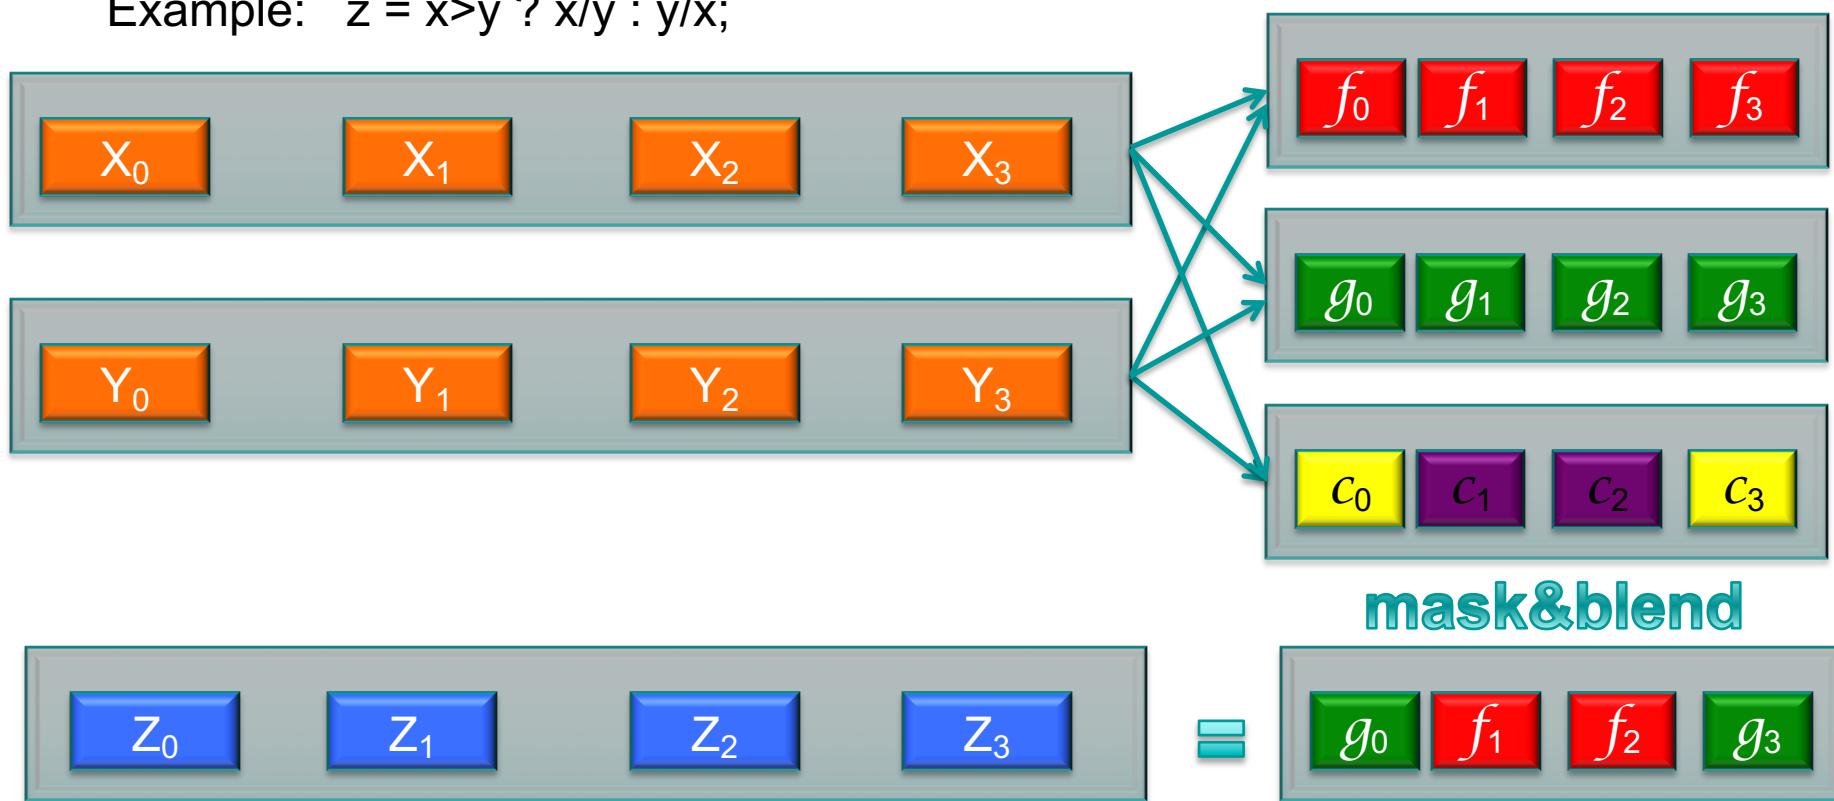
.L4:

```
    movaps b(%rsi), %xmm1  
    addl   $1, %edi  
    addq   $16, %rsi  
    mulps  a-16(%rsi), %xmm1  
    addps  %xmm1, %xmm0  
    cmpl   %edi, %edx  
    ja     .L4  
    haddps %xmm0, %xmm0  
    haddps %xmm0, %xmm0  
    cmpl   %ecx, %eax  
    je     .L15
```

Conditional Code

$$Z_i = c(X_i, Y_i) \ ? \ f(X_i, Y_i) \ : \ g(X_i, Y_i)$$

Example: $z = x > y ? x/y : y/x;$



Conditional code : <https://godbolt.org/g/shCUuE>

```
double x[1024],y[1024] ,z[1024]; int N=1024;
void foo() {
    for (int i=0; i!=N; i++) {
        if (x[i]>y[i]) z[i]=x[i]/y[i];
        else z[i]=y[i]/x[i];
    } // exercise: write it with only one division
}
```

loop in vector notation:

```
for (i=0; i<N/4; ++i){
    V4 vx = &((V4 const *)x[i]);
    V4 vy = &((V4 const *)y[i]);
    V4 & vz = *((V4 *)z[i]);
    I4 t = vx > vy;
    // compute both branches
    auto a = vx / vy; auto b = vy / vx;
    // mask and "blend"
    vz = t ? a : b;
```

```
foo():
.L2:
    xorl    %eax, %eax
    movapd  x(%rax), %xmm2
    addq    $16, %rax
    movapd  y-16(%rax),
%xmm0
    movapd  %xmm2, %xmm3
    divpd   %xmm0, %xmm3
    movapd  %xmm0, %xmm1
    cmplpd  %xmm2, %xmm0
    divpd   %xmm2, %xmm1
    blendvpd %xmm0, %xmm3,
%xmm1
    movaps  %xmm1, z-
16(%rax)
    cmpq    $8192, %rax
    jne     .L2
    ret
```

} **The compiler is able to make the transformation of a condition in “compute, mask and blend” if code is not too complex**

Vectorization of “math function”

- Currently linux libm does not support vectorization
- Exploit compiler + vendor libraries
 - Requires licensed libs by intel and/or amd
 - No change in user code: just recompile
 - -mveclibabi=svml -L/.../lib/intel64 -lsvml -lirc
 - -mveclibabi=acml -L/.../lib/amd64 -lacml -lamdlibm
- Exploit auto-vectorization
 - Modified cephes library (or other open source)
 - Requires header-file + different function names
 - Look into vdt/ (<https://svnweb.cern.ch/trac/vdt>)
 - Project for a student: port vdt to gcc simd-vectors
 - For your exercise use SimpleSinCos.h
 - simpleSin(x), simpleCos(x), simpleSinCos(x,s,c)

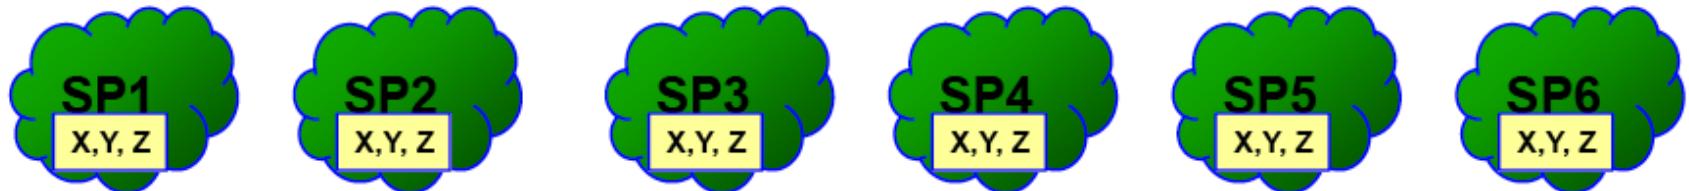
[https://github.com/CppCon/CppCon2014/tree/master/Presentations
/Data-Oriented%20Design%20and%20C%2B%2B](https://github.com/CppCon/CppCon2014/tree/master/Presentations/Data-Oriented%20Design%20and%20C%2B%2B)

<http://aras-p.info/texts/files/2018Academy%20-%20ECS-DoD.pdf>

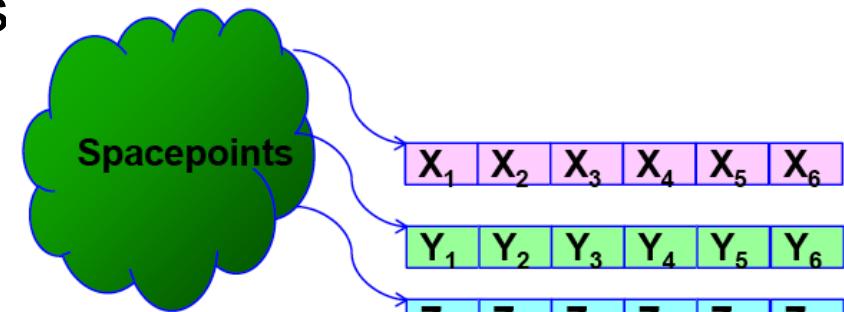
DATA ORGANIZATION

Data Organization: AoS vs SoA

- Traditional Object organization is an Array of Structures
 - Abstraction often used to hide implementation details at object level



- Difficult to fit stream computing
- Better to use a Structure of Arrays
 - (column-wise storage)
- OO can wrap SoA as the AoS
 - Move abstraction higher
 - Expose data layout to the compiler
- Explicit copy in many cases more efficient
 - (notebooks vs whiteboard)



AoS

VS

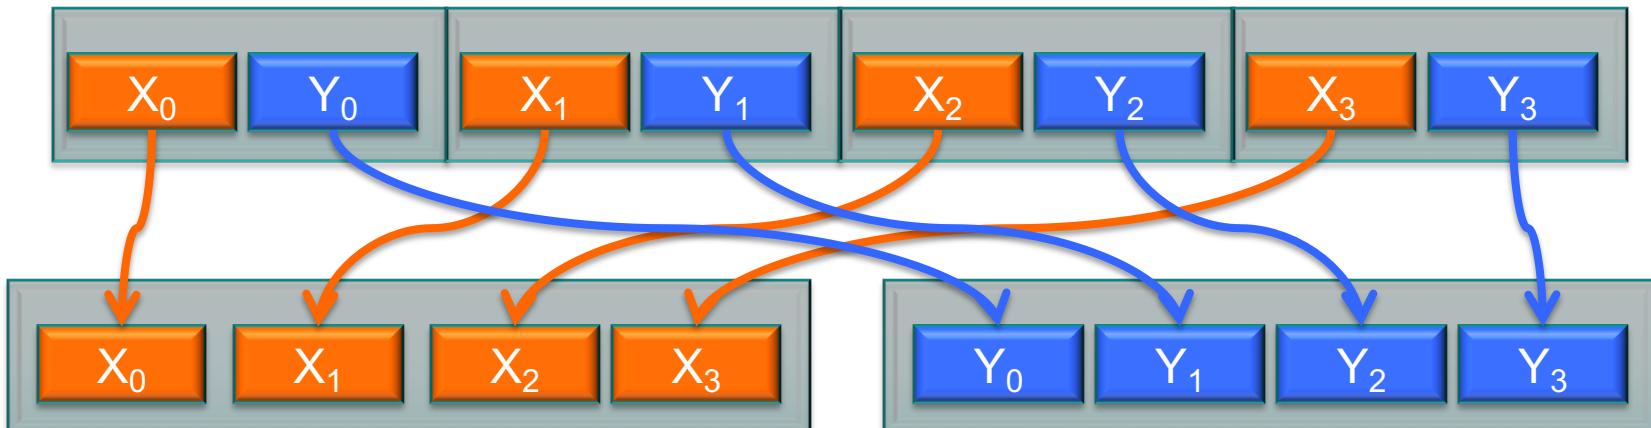
SoA

```
struct Point {  
    float x,y,z;  
};  
  
struct Points {  
    std::vector<Point> p;  
};
```



```
struct Points {  
    std::vector<float> x,y,z;  
};
```

The choice of AoS vs SoA depends on the access patterns in the application
Operations on AoS can be vectorized at the cost of *shuffles* or *permutations*
For large structures this becomes not profitable pretty fast.
In general SoA shall be preferred to fully exploit vectorization

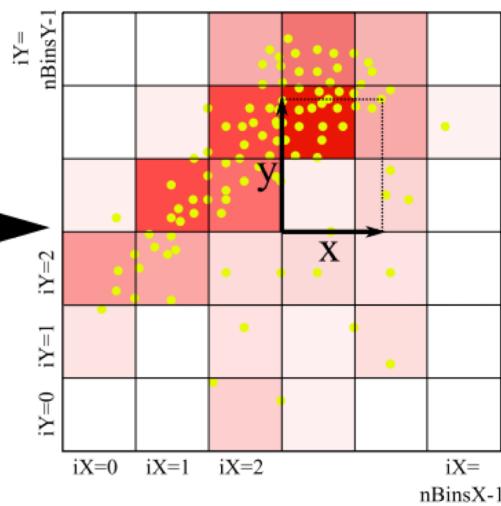
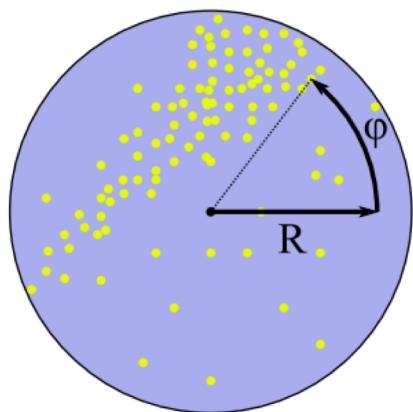


Don't stop the stream!

- Vector code is effective as long as
 - We do not go back to memory
 - Operate on local registries as long as possible
 - We maximize the number of useful operations per cycle
 - Conditional code is a killer!
 - Better to compute all branches and then blend
- Algorithms optimized for sequential code are not necessarily still the fastest in vector
 - Often slower (older...) algorithms perform better

Excercise

<https://infn-esc.github.io/esc16/vectorization/binning.html>



Take points in polar coordinates
Fill a two dimensional
cartesian Histogram

```
struct Point { float phi,r; };
struct Points {
    std::array<Point,N> p;
};
Points points;
constexpr int NBin = 100;
struct Hist {
    int bin[NBin+1][NBin+1] = {0};
};
float binwidth = 2./NBin;
Hist h;
```

```
For(auto const & p : points.p) {
    auto x = p.r*std::cos(p.phi);
    auto y = p.r*std::sin(p.phi);
    int xbin = (x+1.f)/binwidth;
    int ybin = (y+1.f)/binwidth;
    assert(xbin>=0 && ybin>=0);
    assert(xbin<101 && ybin<101);
    ++h.bin[xbin][ybin];
}
```

Exercise 1

- Open SimpleVectorization.cc
<http://goo.gl/SAiici>
- Compile it with O3, Ofast, more fancy options
 - Analyze the compiler report
 - Correlate with code and generated instruction
 - Is vectorizing everything?
 - (if too complex, comment-out all but one function)
 - Add “-fopt-info-vec”
 - Try “-fopt-info-vec-missed -fno-tree-slp-vectorize”
objdump -S -r -C --no-show-raw-instr –w xyz.o | less
build SimpleVectorization_vect.s;
cat SimpleVectorization_vect.s| less

Exercise 2

<https://infn-esc.github.io/esc18/vectorization/pi.html>

- Open pi.cpp <http://goo.gl/zI3WqS>
 - Compile it with -O2, -Ofast, more fancy options
 - Change double in float
 - Try to vectorize using native vectors...

Pi Program: Vectorization with intrinsics (SSE)



See <https://software.intel.com/sites/landingpage/IntrinsicsGuide/>

```
float pi_sse(int num_steps)
{ float scalar_one =1.0, scalar_zero = 0.0, ival, scalar_four =4.0, step, pi, vsum[4];
  step = 1.0/(float) num_steps;

  __m128 ramp  = _mm_setr_ps(0.5, 1.5, 2.5, 3.5);
  __m128 one   = _mm_load1_ps(&scalar_one);
  __m128 four   = _mm_load1_ps(&scalar_four);
  __m128 vstep  = _mm_load1_ps(&step);
  __m128 sum    = _mm_load1_ps(&scalar_zero);
  __m128 xvec;  __m128 denom; __m128 eye;

  for (int i=0;i< num_steps; i=i+4){      // unroll loop 4 times
    ival    = (float)i;                  // and assume num_steps%4 = 0
    eye    = _mm_load1_ps(&ival);
    xvec   = _mm_mul_ps(_mm_add_ps(eye,ramp),vstep);
    denom  = _mm_add_ps(_mm_mul_ps(xvec,xvec),one);
    sum    = _mm_add_ps(_mm_div_ps(four,denom),sum);
  }
  _mm_store_ps(&vsum[0],sum);
  pi = step * (vsum[0]+vsum[1]+vsum[2]+vsum[3]);
  return pi;
} // credit T.M.
```